

Programming Assignment #3: Text Prediction with Tries

Fall 2020

Due: Sunday, November 22, *before* 11:59 PM

Abstract

In this programming assignment, you will gain experience with an advanced tree data structure that is used to store strings, and which allows for efficient insertion and lookup: a trie! You will then use this data structure to complete a text prediction task.

In this assignment, you will also learn how to use valgrind to test your programs for memory leaks.

By completing this assignment and reflecting upon the awesomeness of tries, you will fortify your knowledge of algorithms and data structures and solidify your mastery of many C programming topics you have been practicing all semester: dynamic memory management, file I/O, processing command line arguments, dealing with structs and pointers to structs, and so much more.

In the end, you will have implemented a tremendously useful data structure that has many applications in text processing and corpus linguistics.

Attachments

TriePrediction.h, printTrie.txt,
corpus{01-09}.txt, input{01-09}.txt, output{01-09}.txt

Deliverables

TriePrediction.c

(Note: Capitalization of your filename matters!)
(This project has been adapted from Sean Szumlanski's projects)

1. Overview: Tries

We have seen in class that the trie data structure can be used to store strings. It provides for efficient string insertion and lookup; insertion into a trie is $O(k)$ (where k is the length of the string being inserted), and searching for a string is an $O(k)$ operation (worst-case). In a trie, a node does not store the string it represents; rather, the *edges* taken to reach that node from the root indicate the string it represents. Each node contains a *flag* (or *count*) variable that indicates whether the string represented by that node has been inserted into the trie (or *how many times* it has been inserted). For example:

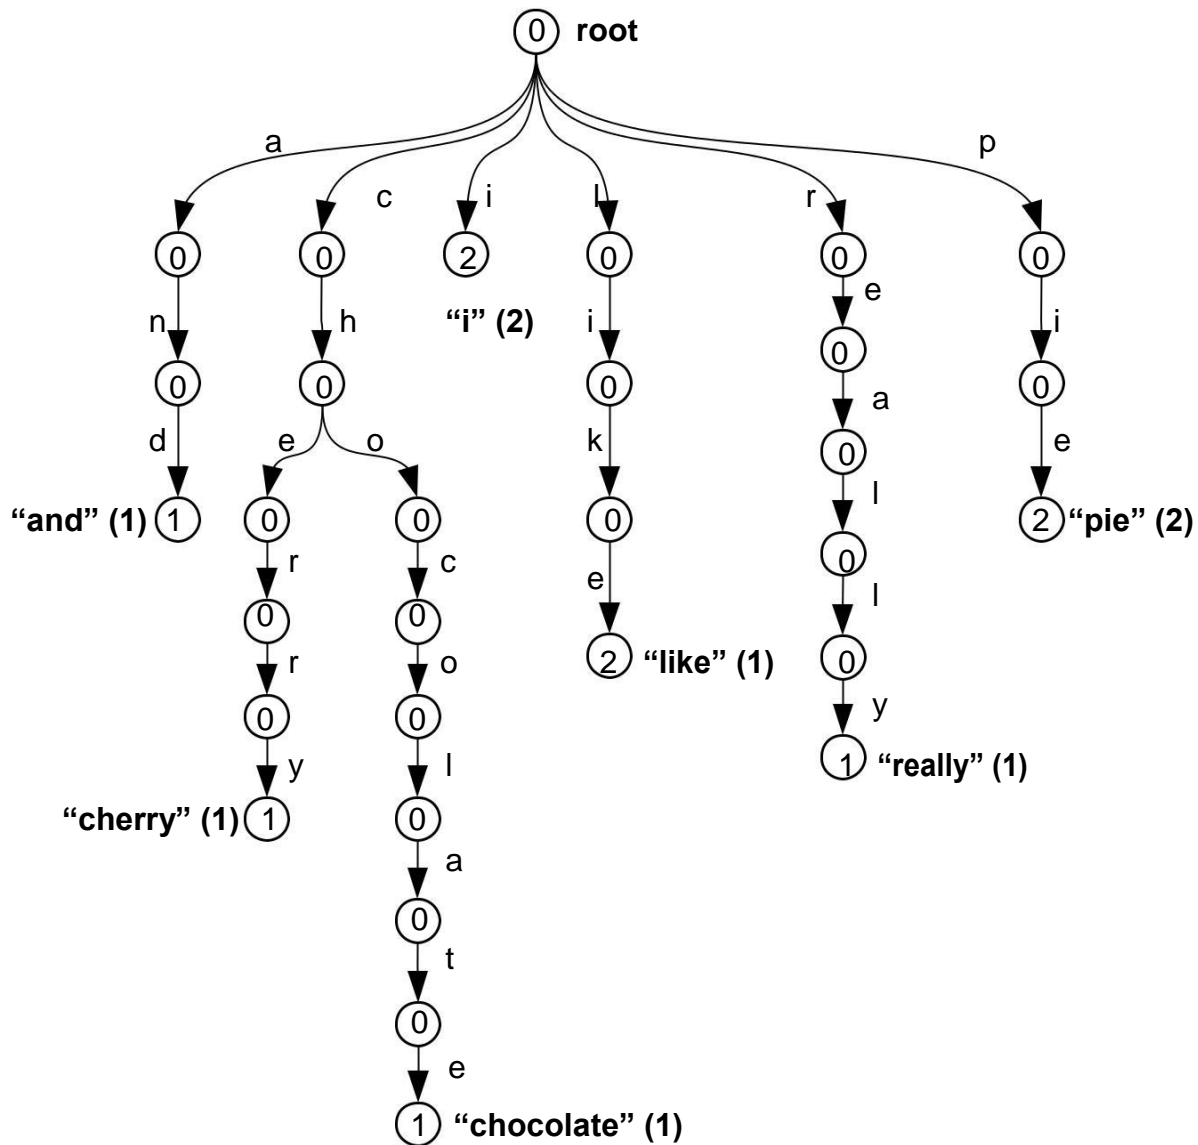


Figure 1:

This is a trie that codifies the words "and," "cherry," "chocolate," "i," "like," "really," and "pie." The strings "i," "like," and "pie" are represented (or counted) twice. All other strings are counted only once.

1.1. TrieNode Struct (TriePrediction.h)

In this assignment, you will insert words from a *corpus* (that is, a body of text from an input file) into a trie. The struct you will use for your trie nodes is as follows:

```
typedef struct TrieNode
{
    //    number of times this string occurs in the corpus int count;

    //    26 TrieNode pointers, one for each letter of the alphabet struct TrieNode
    *children[26];

    //    the co-occurrence subtrie for this string
    struct TrieNode *subtrie;
} TrieNode;
```

You must use this trie node struct, which is specified in TriePrediction.h without any modifications. You should `#include` the header file from TriePrediction.c like so:

```
#include "TriePrediction.h"
```

Notice that the trie node, because it only has 26 children, represents strings in a case insensitive way (i.e., “apple” and “AppLE” are treated the same in this trie).

1.2. Subtries: Contextual Co-occurrence and Predictive Text

Words that appear together in the same sentence are said to “co-occur.” In this program, we’ll be interested in contextual co-occurrence and predictive text – namely, given some word, w , what are all the words that we see immediately after w in the sentences in some corpus? Consider, for example, the following sentence:

I like cherry pie, and I really like chocolate pie.

In the example sentence, the word “I” is followed by the words “like” and “really”, the word “pie” is followed by the word “and”, and so on.

To track co-occurrence, for each word w , we’ll place each word that directly follows w into the subtrie of w . For example, if we place these terms (and their associated counts) into their own tries, those tries will look like this:

(See following page for diagram.)




Figure 2:
Co-occurrence subtries for “I” (left) and “like” (right), based on the sentence, “I like cherry pie, and I really like chocolate pie.”

In Figure 2, the trie on the left is what we will call the *co-occurrence subtrie* for the word “I,” based on the example sentence above (*I like cherry pie, and I really like chocolate pie*). Its root should be stored in the *subtrie* pointer field of the node marked “I” in the original trie diagram in Figure 1 (see page 2 above).

Similarly, the trie on the right in Figure 2 is the co-occurrence subtrie for the word “like.” Its root should be stored in the *subtrie* pointer field of the node marked “like” in the original trie diagram in Figure 1 (see page 2, above).

Within these subtries, all *subtrie* pointers should be initialized to NULL, because we will NOT produce sub-subtries in this assignment!

2. Input Files and Output Format

2.1. Command Line Arguments

Your program will take two command line arguments, specifying two files to be read at runtime:

```
./a.out corpus01.txt input01.txt
```

The first filename specifies a corpus that will be used to construct your trie and subtries. The second filename specifies an input file with strings to be processed based on the contents of your trie.

2.2. Corpus File

2.2.1 Corpus File Format

The corpus file contains a series of sentences. Each line contains a single sentence with at least 1 word and no more than 30 words. Each word contains fewer than 1024 characters, some of which may be punctuation characters.

Each sentence will contain exactly one of the following punctuators, and it will occur at the very end of the sentence: period ('.'), exclamation point ('!') or question mark ('?'). Other punctuators may occur throughout the sentence, including (but not limited to) commas (','), colons (':'), semicolons (';'), apostrophes ("" and ""), quotation marks ("" and ""), and so on.

All words will have at least one alphabetic character ('a' through 'z' or 'A' through 'Z'). Words will *not* contain any numeric characters ('0' through '9'). All punctuators should be stripped away from all words before entering them into the trie. So, for example, the word "don't" will be entered into the trie as "dont".

For example:

corpus01.txt:

```
I like cherry pie, and I really like chocolate pie.
```

corpus02.txt:

```
'tweren't my fault, I swear!  
And now we're on to the second line of text.
```

2.2.2 Building the Main Trie

First and foremost, each word from the corpus file should be inserted into your trie. If a word occurs multiple times in the corpus, you should increment *count* variables accordingly. For example, the trie in Figure 1 (see page 2, above) corresponds to the text given in the corpus01.txt file above.

2.2.3 Building Subtries

For each sentence in the corpus, update the co-occurrence subtrie for each word in that sentence. The structure of the co-occurrence subtries is described above in Section 1.2, “Subtries: Contextual Co-occurrence and Predictive Text.”

If a string in the main trie is not followed by any other words in the corpus, its subtrie pointer should be NULL.

2.3. Input File

The input file (the second filename specified as a command line argument) will contain any number of lines of text. Each line will correspond to one of the following three commands:

Command: <string>

Description:

<string> is a string containing fewer than 1024 characters. The string will contain alphabetical characters only ('A' through 'Z' and/or 'a' through 'z'), and it appears on a line by itself.

If <string> is in the trie, you should print the string, followed by a printout of its subtrie contents, using the output format shown in the sample outputs for this program. (Note that when printing a subtrie, the words in the subtrie are preceded by hyphens.)

If <string> is in the trie, but its subtrie is empty, you should print that string, followed on the next line by “(EMPTY)”.

If <string> is not in the trie at all, you should print that string, followed on the next line by “(INVALID STRING)”.

Command: !

Description:

The character '!' will appear on a line by itself. When you encounter this command, you should print the trie using the output format shown in the sample outputs for this

program. (When printing the main trie, there are no hyphens preceding the words on each line. See sample output files, or see the sample output below on pg. 6.)

Command: @ <string> <n>

Description:

This is the text prediction command.

<string> is a string containing fewer than 1024 characters. The string will contain alphabetical characters only ('A' through 'Z' and 'a' through 'z'). <n> is an integer.

When you encounter this command, you should print the following sequence of $(n + 1)$ words:

$$w_0 \ w_1 \ w_2 \ \dots \ w_n$$

where w_0 is <string>, and for $1 \leq i \leq n$, w_i is the word that most frequently follows word w_{i-1} in the corpus. Note that the words are separated by spaces, and there is never a space after the last word on one of these lines of output.

Furthermore, w_0 is always capitalized in the output exactly as it was capitalized in the input file, but words w_1 through w_n should appear in all lowercase.

If <string> does not appear in the trie, it should appear on this line of output by itself. If some w_i does not have any words in its subtrie, the sequence should terminate prematurely. Again, this line of output should *not* have a trailing space at the end of it, even if it terminates prematurely.

2.4. Sample Input and Output Files

For example, the following corpus and input files would produce the following output:

corpus03.txt:

```
I like cherry pie and chocolate pie.
```

input03.txt:

```
!  
chocolaTE  
apricot  
@ I 11  
@ chocolate 1  
@ persimmon 20
```

(Continued on the following page.)

output03.txt:

```
and (1)
cherry (1)
chocolate (1)
i (1) like
(1) pie (2)
chocolaTE -
pie (1)
apricot

(INVALID STRING)
I like cherry pie and chocolate pie and chocolate pie and chocolate chocolate pie

persimmon
```

Note that a word might be in the trie but have an empty subtrie. Consider the following example:

corpus04.txt:

```
Spin straw to gold.
Spin all night long.
Spin spin spin.
Spindle.
```

input04.txt:

```
spin
spindle
nikstlitslepmur
```

output04.txt:

```
spin
- all (1)
- spin (2)
- straw (1) spindle
(EMPTY)
nikstlitslepmur

(INVALID STRING)
```

You must follow the output format above precisely. Be sure to consult the included text files for further examples.

I have included some functions that will help you print the contents of your trie(s) in the required format, because I think those functions are a bit too tricky to expect you to write them on your own. See `printTrie.txt` (attachment) for those functions. You're welcome to copy and paste them into your `TriePrediction.c` source file if you want to use them, or you can write your own.

Also, studying and understanding those functions will serve as a launching point for you to write the other functions required to get this program working.

2.5. Final Thoughts on Test Cases, Input, and Output

I have attached several sample input and output files so you can check that your program is working as intended. Be sure to use diff to make sure your output matches ours exactly. I also encourage you to develop your own test cases.

3. Special Requirement: Memory Leaks and Valgrind

Part of the credit for this assignment will be awarded based on your ability to implement the program without any memory leaks. To test for memory leaks, you can use a program called valgrind.

To run your program through valgrind at the command line, compile your code with the -g flag, and then run valgrind, like so:

```
gcc TriePrediction.c -g
valgrind --leak-check=yes ./a.out
```

For help deciphering the output of valgrind, see: <http://valgrind.org/docs/manual/quick-start.html>

Note that if you do not use fclose() to explicitly close all open files before your program terminates, valgrind might alert you that your program has a memory leak.

In the output of valgrind, the magic phrase you're looking for to indicate that you have no memory leaks is:

All heap blocks were freed – no leaks are possible

(Continued on the following page.)

4. Function Requirements

4.1. Required Function

You have a lot of leeway with how to approach this assignment. There are only five required functions. How you structure the rest of your program is up to you.

```
int main(int argc, char **argv);
```

Description: You must write a main() function.

Returns: Your main() must return zero.

```
TrieNode *buildTrie(char *filename);
```

Description: filename is the name of a corpus text file to process. Open the file and create a trie (including all its appropriate subtries) as described above.

Returns: The root of the new trie.

```
TrieNode *destroyTrie(TrieNode *root);
```

Description: Free all dynamically allocated memory associated with this trie.

Returns: NULL

```
double difficultyRating(void);
```

Returns: A double indicating how difficult you found this assignment on a scale of 1.0 (ridiculously easy) through 5.0 (insanely difficult).

```
double hoursSpent(void);
```

Returns: An estimate (greater than zero) of the number of hours you spent on this assignment.

(Continued on the following page.)

4.2. Suggested Functions

These functions are not required, but I think they will simplify your task immensely if you implement them properly and call them when processing your corpus/input files. Think of these function descriptions as hints at how to proceed. If you want, you can even implement these functions with different parameters, return types, and so on.

```
TrieNode *createTrieNode(void);
```

Description: Dynamically allocate space for a new TrieNode struct. Initialize all the struct members appropriately.

Returns: A pointer to the new node.

```
TrieNode *getNode(TrieNode *root, char *str);
```

Description: Searches the trie for the specified string, str.

Returns: If the string is represented in the trie, return its terminal node (the last node in the sequence, which represents that string). Otherwise, return NULL.

```
void insertString(TrieNode *root, char *str);
```

Description: Inserts the string str into the trie. Since it has no return value, it assumes the root already exists (i.e., root is not NULL). If str is already represented in the trie, simply increment its count member.

```
void mostFrequentWord(TrieNode *root, char *str);
```

Description: Searches the trie for the most frequently occurring word and copies it into the string variable passed to the function, str. str should be initialized before calling this function, and it should be long enough to hold the string that will be written to it. If there are multiple strings in the trie that are tied for the most frequently occurring, populate str with the one that comes first in alphabetical order. If the trie is empty, set str to the empty string ("").

Hint: You might find it easier to write helper functions that you call from within this function.

```
void stripPunctuators(char *str);
```

Description: Takes a string, str, and removes all punctuation from the string. For example, if str contains the string "Hello!" when the function is called, then str should contain the string "Hello" when the function returns. When writing this function, you might find C's built-in isalpha() function (from ctype.h) to be helpful.

5. Compilation and Testing (Linux/Mac Command Line)

To compile at the command line:

```
gcc TriePrediction.c
```

By default, this will produce an executable file called `a.out` that you can run by typing, e.g.:

```
./a.out corpus01.txt input01.txt
```

If you want to name the executable something else, use:

```
gcc TriePrediction.c -o TriePrediction.exe
```

and then run the program using:

```
./TriePrediction.exe corpus01.txt input01.txt
```

Running the program could potentially dump a lot of output to the screen. If you want to redirect your output to a text file in Linux, it's easy. Just run the program using the following:

```
./TriePrediction.exe corpus01.txt input01.txt > whatever.txt
```

This will create a file called `whatever.txt` that contains the output from your program.

Linux has a helpful command called `diff` for comparing the contents of two files, which is really helpful here since we've provided sample output files. You can see whether your output matches ours exactly by typing, e.g.:

```
diff whatever.txt output01.txt
```

If the contents of `whatever.txt` and `output01.txt` are exactly the same, `diff` won't have any output. It will just look like this:

```
navid@ubuntu:~$ diff whatever.txt output01.txt
navid@ubuntu:~$ _
```

If the files differ, it will spit out some information about the lines that aren't the same. For example:

```
navid@ubuntu:~$ diff whatever.txt output01.txt 12c12
< - chocolate (1)
---
> - cherry (1)
navid@ubuntu:~$ _
```

6. Grading Criteria and Submission

6.1. Deliverables

Submit a single source file, named `TriePrediction.c`, via CANVAS. The source file should contain definitions for the required function (listed above), as well as any auxiliary functions you decide to implement. Don't forget to `#include "TriePrediction.h"` in your source code. Your program must compile with the following:

```
gcc TriePrediction.c
```

Be sure to include your name and ID in a comment at the top of your source file.

6.2. Additional Restrictions: Use Tries; Do Not Use Global Variables

You must use tries to receive credit for this assignment. Also, please do not use global variables in this program. Doing so may result in a huge loss of points.

6.3. Grading

The expected scoring breakdown for this programming assignment is:

60%	Correct Output for Test Cases
20%	Unit Testing of <code>buildTrie()</code> Function
10%	No Memory Leaks; Passes <code>valgrind</code> Tests
10%	Comments and Whitespace

Your program must compile and run to receive credit. Programs that do not compile will receive an automatic zero.

Your grade will be based primarily on your program's ability to compile and produce the *exact* output expected. Even minor deviations (such as capitalization, punctuation, or whitespace errors) in your output will cause your program's output to be marked as incorrect, resulting in severe point deductions. The same is true of how you name your functions and their parameters. Please be sure to follow all requirements carefully and test your program thoroughly.

For this program, we will also be unit testing your code. That means we will devise tests that determine not only whether your program's overall output is correct, but also whether your `buildTrie()` function does exactly what it is required to do. So, for example, if your program produces correct output but your `buildTrie()` function is simply a skeleton that returns `NULL` no matter what parameters you pass to it, or if it produces a totally funky, malformed trie, your program will fail the unit tests.

6.4. Closing Remarks

Start early. Work hard. Ask questions. Good luck!